

# Finding and Reproducing Heisenbugs in Concurrent Programs

Madanlal Musuvathi   Shaz Qadeer   Thomas Ball   Gerard  
Basler

Microsoft Research

OSDI '08

Presentation given by Wang Yuanxuan

# Outline

- 1 Introduction
- 2 Example
- 3 The CHESS scheduler
- 4 Exploring nondeterminism
- 5 Evaluation
- 6 Conclusions

# Outline

- 1 Introduction
  - Introduction to CHESS
  - Architecture
- 2 Example
- 3 The CHESS scheduler
- 4 Exploring nondeterminism
- 5 Evaluation
- 6 Conclusions

# Heisenbugs

## Definition

Subtle interactions among threads and the timing of asynchronous events can result in concurrency errors that are hard to find, reproduce and debug. Stories are legend of so-called "**Heisenbugs**" that occasionally surface in systems that have otherwise been running reliably for months.

# CHES

## Definition

A tool for systematic and deterministic testing of concurrent programs.

## Features

CHES takes complete control over the scheduling of threads and asynchronous events.

- Capability to reproduce the erroneous thread interleaving.
- Systematic enumeration techniques to force every run of the program along a different threading interleaving.

# CHESS

## Definition


A tool for systematic and deterministic testing of concurrent programs.

## Challenges

Challenges to build a systematic testing tool for concurrent programs

- Avoid perturbing the system under test.
- Accomplish the nontrivial task of capturing and exploring all interleaving nondeterminism.
- Explore the space of thread interleavings intelligently.

# CHESSE Architecture



# Outline

- 1 Introduction
- 2 **Example**
  - Problem
  - Solution
- 3 The CHESS scheduler
- 4 Exploring nondeterminism
- 5 Evaluation
- 6 Conclusions



# Example

## Example

How CHESS was used to reproduce a Heisenbug in CCR, a .NET library for efficient asynchronous concurrent programming.

## The Bug

The entire test run consists of many smaller unit concurrency tests. The failing test(which did not terminate) previously had not failed for many months.

# Steps

- Changed the harness so it ran the test just once.

# Steps

- Changed the harness so it ran the test just once.
- Ran the test under CHESS

# Steps

- Changed the harness so it ran the test just once.
- Ran the test under CHES
- In just over 20secs, CHES reported a deadlock after exploring 6737 different thread interleavings.

# Steps

- Changed the harness so it ran the test just once.
- Ran the test under CHES
- In just over 20secs, CHES reported a deadlock after exploring 6737 different thread interleavings.
- Let CHES reproduce the last deadlock scenario.

# Steps

- Changed the harness so it ran the test just once.
- Ran the test under CHES
- In just over 20secs, CHES reported a deadlock after exploring 6737 different thread interleavings.
- Let CHES reproduce the last deadlock scenario.
- Ran CHES on the offending schedule under the control of a standard debugger.

# Steps

- Changed the harness so it ran the test just once.
- Ran the test under CHES
- In just over 20secs, CHES reported a deadlock after exploring 6737 different thread interleavings.
- Let CHES reproduce the last deadlock scenario.
- Ran CHES on the offending schedule under the control of a standard debugger.
- The source of the bug was identified.

# Outline

- 1 Introduction
- 2 Example
- 3 The CHES scheduler
  - Handling input nondeterminism
  - Choosing the right abstraction layer
  - The happens-before graph
  - Capturing the happens-before graph
  - Capturing data-races by single-threaded execution
- 4 Exploring nondeterminism
- 5 Evaluation



# Primary Goal

- Capture all the nondeterminism during a program execution.
- Expose these nondeterministic choices to a search engine

# Handling input nondeterminism

- Shift the onus of generating deterministic inputs to the user.
- Consider an elaborate log and replay mechanism unnecessary.
- Log and replay input values that are not easily controlled by the user.

# A Wrapper Library


- The scheduler redirects calls to concurrency primitives.
- By including complex primitives as part of the program, the scheduler only needs to understand the simpler primitives.

# The happens-before graph

## Definition

The graph capturing the relative execution order of the threads in a concurrent execution.

## Example



# The happens-before graph

## Definition

The graph capturing the relative execution order of the threads in a concurrent execution.

## Benefits

- Provide a common framework for reasoning about all the different synchronization primitives used by a program.
- Abstract the timing of instructions in the execution.

# The happens-before graph

## Definition

The graph capturing the relative execution order of the threads in a concurrent execution.

## Node

Each node is annotated with a triple  
(task, synchronization variable, operation)

# The happens-before graph

## Definition

The graph capturing the relative execution order of the threads in a concurrent execution.

## Tasks

- Threading executing an instruction
- Threadpool work items
- asynchronous callbacks
- timer callbacks

# The happens-before graph

## Definition

The graph capturing the relative execution order of the threads in a concurrent execution.

## Operations

CHES only needs to understand

- isWrite. if true, two sets of edges are created)
- isRelease. Needed by the search module.



# Wrapper library

- Determine whether a task may be disabled by executing a potentially blocking API call.
- Label each call to the API with an appropriate triple.
- Inform the CHESSE scheduler about the creating and termination of a task.

# Robustness in design

- Conservatively setting isWrite to true only adds extra edges in the happens-before graph
- Conservatively setting isRelease to true might creating some wasteful work.

# Code complexity

API	No. of wrappers	LOC
Win32	134	2512
.NET	64	1270
Singularity	37	876

# Data-race problem

## Problem

Most concurrent programs contain data-races, which cannot be captured by the wrappers.

## Solutions

Enforce single-threaded execution to ensure that two threads cannot concurrently access memory locations. All data-races occur in the order in which CHESS schedules the threads.

# Data-race problem

## Problem

Most concurrent programs contain data-races, which cannot be captured by the wrappers.

## Downsides

- Slow down the execution of the program. But this lost performance can be recovered by running multiple CHES instances in parallel.
- CHES may not be able to explore both of the possible outcomes of a data-race. Can be addressed by running a data-race detector.

# Outline

- 1 Introduction
- 2 Example
- 3 The CHESS scheduler
- 4 Exploring nondeterminism
  - Basic search operation
  - Dealing with imperfect replay
  - Tackling state-space explosion
  - Monitoring executions
- 5 Evaluation

# Basic search operation

## Strategy

- CHES repeatedly executes the same test driving each iteration of the test through a different schedule

# Basic search operation

## Strategy

- CHESS repeatedly executes the same test driving each iteration of the test through a different schedule
- In each iteration, the scheduler works in three phases: replay, record, and search.



# Basic search operation

## Strategy

- CHES repeatedly executes the same test driving each iteration of the test through a different schedule
- In each iteration, the scheduler works in three phases: replay, record, and search.
- Replay: replay a sequence of scheduling choices from a trace file.

# Basic search operation

## Strategy

- CHES repeatedly executes the same test driving each iteration of the test through a different schedule
- In each iteration, the scheduler works in three phases: replay, record, and search.
- Replay: replay a sequence of scheduling choices from a trace file.
- Record: the scheduler behaves as a fair, non-preemptive scheduler. It also record the thread scheduled at each schedule point with the set of threads enabled at each point.

# Basic search operation

## Strategy

- CHES repeatedly executes the same test driving each iteration of the test through a different schedule
- In each iteration, the scheduler works in three phases: replay, record, and search.
- Replay: replay a sequence of scheduling choices from a trace file.
- Record: the scheduler behaves as a fair, non-preemptive scheduler. It also record the thread scheduled at each schedule point with the set of threads enabled at each point.
- Search: determine the schedule for the next iteration.

# Fail to replay a trace

## Two cases

- The thread to schedule at a scheduling point is disabled.
- A scheduled thread performs a different sequence of synchronization operations than the one present in the trace.

## Solutions

- Lay-initialization
- Interference from environment
- Nondeterministic call
  - `random()`
  - `gettimeofday()`

# State-space explosion

## Problem

Given a program with  $n$  threads that execute  $k$  atomic steps in total, it is very easy to show that the number thread interleavings grows astronomically as  $n^k$ .

# Inserting preemptions prudently

## Reference

Given a program with  $n$  threads that execute  $k$  steps in total, the number of interleavings with  $c$  preemptions grows with  $k^c$ .

## Optimizations

Scope preemptions to code regions of interest, essentially reducing  $k$ .

- A significant portion of the synchronization operations occur in system functions.
- A large number of the synchronizations are due to accesses to volatile variables.

# Monitoring

## Failure catching

- Null dereferences, segmentation faults, crashes.
- User can attach other monitors such as memory-leak detectors.
- Deadlock report – whenever the set of enabled tasks becomes empty.
- Livelock report – user sets an abnormally high bound on the length of the execution.

# Outline

- 1 Introduction
- 2 Example
- 3 The CHES scheduler
- 4 Exploring nondeterminism
- 5 Evaluation**
  - Benchmarks
  - Validating CHES against stress-testing
- 6 Conclusions



# Brief description

## Input programs

Programs	LOC	max threads	max synch.	max preemp.
PLINQ	23750	8	23930	2
CDS	6243	3	143	2
STM	20176	2	75	4
TPL	24134	8	31200	2
ConcRT	16494	4	486	3
CCR	9305	3	226	2
Dryad	18093	25	4892	2
Singularity	174601	14	167924	1

# Bugs found by CHES

Programs	Total	Unk/Unk	Kn/Unk	Kn/Kn
PLINQ	1		1	2
CDS	1		1	
STM	2			
TPL	9	9		
ConcRT	4	4		
CCR	2	1	1	
Dryad	7	7		
Singularity	1		1	
Total	27	21	4	2

# Outline

- 1 Introduction
- 2 Example
- 3 The CHESS scheduler
- 4 Exploring nondeterminism
- 5 Evaluation
- 6 Conclusions**

# Conclusions

- CHES, a systematic testing tool for finding and reproducing Heisenbugs in concurrent programs.
- Achieved by carefully exposing, controlling, and searching all interleaving nondeterminism in a concurrent system.

Thank you!